

Backup, Restore, and Recovery

The database is running.
Yesterday's records are gone.

Day 1: recover a small PostgreSQL database
Day 2: connect recovery to the midterm case

An accidental committed deletion

```
BEGIN;  
DELETE FROM metro_support.ticket_events;  
COMMIT;
```

Illustrative incident: the operator intended to remove one test event.
The omitted WHERE condition removes every event.

A restart preserves the committed change. A replica may copy it.
Recovery needs an earlier usable state.

Availability and recovery cover different failures


Mechanism	What it provides	Important limit
High availability	Continued service after some component failures	Can reproduce a harmful committed write
Backup	A retained recoverable state	Must remain accessible and restorable
Restore	Reconstruction into a target	Needs verification before application use
Disaster recovery	A tested service-recovery procedure	Includes people, access, systems, and dependencies

The failure you plan for determines what you must keep and test.

Recovery point and recovery time objectives

RPO: acceptable data-loss window.

RTO: target duration to restore acceptable service.



t_r : newest recoverable state

t_f : failure

t_a : service acceptable

Data-loss window = $t_f - t_r$

Service-recovery duration = $t_a - t_f$

A worked recovery timeline

Time	Event
14:00	Newest usable export represents this state
14:17	Harmful deletion commits
14:21	Operator recognizes the incident
14:29	Restore and verification finish; service is acceptable

Data gap: 17 minutes. Recovery duration: 12 minutes.


An RPO of 5 minutes and an RTO of 10 minutes would both be missed.
Illustrative scenario. Time gaps do not give the number of lost records.

Logical and physical recovery

Approach	What it reconstructs	What it depends on
Logical export	Objects and rows through database commands	Compatible tools, included dependencies, restore time
Physical backup and logs	Database storage and a supported history of changes	Compatible base backup and retained log sequence

Today: pg_dump and pg_restore for one course schema.
CSV alone does not preserve its keys, checks, indexes, or grants.

Supabase project health and retained backups




Healthy describes present availability.

No backups describes the displayed recovery history.

Supabase Free lacks the automatic backups offered on paid plans. Our lab uses a logical dump and a separate local restore.

Redacted course capture: August 25, 2026. Plan documentation checked September 7, 2026.

A small relational database for the restore



One user can request many tickets. One ticket can have many events.

Notebook fixture: 3 users / 3 tickets / 5 events.

Full Metro Support fixture elsewhere: 8 / 12 / 21. Compare the correct baseline.

Python launches tools; psql sends SQL

```
subprocess.run(  
    PG_PREFIX + [PSQL, '-X', '-d', SOURCE_DB,  
                  '-c', 'SELECT count(*) FROM metro_support.tickets;'],  
    check=True, text=True, capture_output=True,  
)
```

Python waits for psql to finish.

psql connects to SOURCE_DB and submits the SELECT.

PostgreSQL evaluates the SQL and returns the count.

In Colab, PG_PREFIX runs the local command as the postgres operating-system user.

The custom archive contains one schema

```
pg_dump --format=custom \  
  --schema=metro_support --no-privileges \  
  --file=metro_support.dump \  
  --dbname="$SOURCE_DB"
```

Shell form of the notebook operation. SOURCE_DB is a database name, not a password-bearing URL. The notebook supplies the unique name.

Custom archive: --no-owner belongs on pg_restore.
Roles and cloud project settings need separate recovery planning.

Archive inspection and a checksum

```
pg_restore --list metro_support.dump
```

Check	What it establishes	What remains untested
Archive contents	Expected object inventory appears	Whether restoration will succeed
File size	Bytes exist	Completeness or correctness
SHA-256 comparison	Whether bytes match a trusted earlier hash	Whether the original content was trustworthy

A dump can execute SQL chosen by its source. Only use your own course archive in this lab.

A separate empty database receives the restore

```
createdb "$RESTORE_DB"  
pg_restore --exit-on-error --single-transaction \  
  --no-owner --no-privileges \  
  --dbname="$RESTORE_DB" metro_support.dump
```

Stop on failure. Keep this small restore transactional.
Review ownership and grants separately before application use.

SOURCE_DB and RESTORE_DB are different databases on one temporary server.
They still share the risk of losing the runtime.

The source baseline gives specific expected results

Check	Source	Restored target
users / tickets / ticket_events	3 / 3 / 5	3 / 3 / 5
Tickets with no matching requester	0	0
in_progress / open / resolved	1 / 1 / 1	1 / 1 / 1
Named status constraint	tickets_status_allowed	tickets_status_allowed

The notebook compares complete results, including table names.
Equal counts alone could conceal a changed subject or requester.

A known-ticket query checks identity and values

```
SELECT t.ticket_id, t.status, t.subject,  
       u.display_name AS requester  
FROM metro_support.tickets AS t  
JOIN metro_support.users AS u ON u.user_id = t.requester_id  
WHERE t.ticket_id = 1001;
```

1001, open, Streetlight dark near bus stop, Maya Chen

Changing only the subject leaves the row counts unchanged.
Compare values with the source and the intended original record.

A restored constraint must reject the intended violation

```
BEGIN;  
INSERT INTO metro_support.tickets  
    (ticket_id, requester_id, status, subject)  
VALUES (1099, 101, 'almost_done', 'Constraint restore test');  
ROLLBACK;
```

Expected: SQLSTATE 23514 and tickets_status_allowed

A missing table or disconnected server is a different failure.
The notebook also checks that test ticket 1099 was not retained.

Lab: recover and check one ticket of your own

Run the recovery notebook in Colab.

Adapt its known-ticket query to ticket 1002 or 1003.

Explain what that check adds and what remains untested.

Keep the output and short recovery account in the notebook.
Then remove the practice resources with the final cleanup cell.

Individual work. Brightspace: one completed notebook, no second report.

Day 2: the PostgreSQL operations clinic

A working schema must survive ordinary changes and mistakes.

Review: transactions, blocking, and recovery choices
Work: your individual midterm operations case

A small schema change can be rehearsed in a transaction

```
BEGIN;  
ALTER TABLE metro_support.tickets ADD COLUMN follow_up text;  
SELECT column_name FROM information_schema.columns  
WHERE table_schema = 'metro_support'  
      AND table_name = 'tickets' AND column_name = 'follow_up';  
ROLLBACK;
```

Inside the transaction: follow_up exists.

After ROLLBACK: the same catalog query returns no rows.

Disposable notebook fixture only. A transactional change can still take locks.

Rollback, forward repair, and restore

Situation	Candidate response	Decision constraint
Harmful write remains uncommitted	ROLLBACK the owning transaction	Coordinate with the transaction owner
Small committed mistake with a known correction	Reviewed compensating change	Avoid overwriting valid newer work
Committed loss requiring an earlier state	Separate restore and verified recovery	Reconcile later valid changes before cutover

A database backup does not make every incident a restore incident.

A waiting writer may be a transaction problem

Session	Observed state in the Week 5 exercise
A	Changed a ticket and left its transaction open
B	Waiting to update the same ticket
Diagnostic connection	pg_blocking_pids identifies A as B's blocker

After the owner resolves A, B can finish.
Verify the final row and close the practice connections.

This episode does not require restoring yesterday's database.

An improvement starts with a specific operating problem

Problem	Possible improvement	A useful comparison
Reader can write data	Least-privilege role	Allowed SELECT and denied write
User can see another user's rows	Row-level access policy	Two identities with different visible rows
Small queue requires a broad scan	Query-appropriate index	Same result IDs, different plan work
Application needs an additional field	Safe schema change	Old behavior still works after the change

The midterm offers these alternatives. Choose one relevant to your case.

A recovery plan identifies the target and acceptance checks

Purpose: recover the Metro Support schema after accidental loss.

Target: an approved empty database, separate from the source.

Checks: expected objects and counts, a known ticket, a report, and the intended constraint failure.

Before application use: restore reviewed access and verify a safe client path.

Midterm: a recovery plan is required. An executed restore is optional extra work. Keep planned checks distinct from checks you actually ran.

A concise handoff explains the recovery result

I restored the three-table practice archive into a separate database. The counts and grouped status report matched the source. Ticket 1001 retained its subject and requester. The restored status rule rejected the invalid test value. I have not restored application roles or tested application reconnection, so this drill does not establish that the service is ready for traffic.

Worked example from the notebook. Use your own results and limitations when explaining your midterm case.

Clinic resources and the next step

Chapter 8: recovery commands and the runbook example

Week 5 notebook: the blocking experiment

Canonical midterm assignment: requirements and rubric

Use the remaining class work to repair one incomplete part, rerun it, and update the explanation beside the result.

Next week: NoSQL models and JSON, compared with the relational strengths you now know.